

SGD Algorithm to predict movie ratings

1. Download the data from [here](https://drive.google.com/open?id=1-1z7iDB52cB6_Jp07Dqa-eOYSs-mivpq) (https://drive.google.com/open?id=1-1z7iDB52cB6_Jp07Dqa-eOYSs-mivpq).
2. the data will be of this formate, each data point is represented as a triplet of user_id, movie_id and rating

user_id	movie_id	rating
77	236	3
471	208	5
641	401	4
31	298	4
58	504	5
235	727	5

task 1: Predict the rating for a given (user_id, movie_id) pair

- μ : scalar mean rating
- b_i : scalar bias term for user i
- c_j : scalar bias term for movie j
- u_i : K-dimensional vector for user i
- v_j : K-dimensional vector for movie j


then the predicted rating \hat{y}_{ij} for user i , movie j pair is calculated as $\hat{y}_{ij} = \mu + b_i + c_j + u_i^T v_j$ here we will be finding the best values of b_i and c_j using SGD algorithm with the optimization problem for N users and M movies is defined as

$$L = \min_{b, c, \{u_i\}_{i=1}^N, \{v_j\}_{j=1}^M} \alpha \left(\sum_j \sum_k v_{jk}^2 + \sum_i \sum_k u_{ik}^2 + \sum_i b_i^2 + \sum_j c_j^2 \right) + \sum_{i,j \in \mathcal{I}^{\text{train}}} (y_{ij} - \mu - b_i - c_j - u_i^T v_j)^2$$

TASK: 1

SGD Algorithm to minimize the loss

1. for each unique user initialize a bias value B_i randomly, so if we have N users B will be a N dimensional vector, the i^{th} value of the B will corresponds to the bias term for i^{th} user
2. for each unique movie initialize a bias value C_j randomly, so if we have M movies C will be a M dimensional vector, the j^{th} value of the C will corresponds to the bias term for j^{th} movie
3. Construct adjacency matrix with the given data, assumeing its [weighted un-directed bi-partited graph](https://en.wikipedia.org/wiki/Bipartite_graph) (https://en.wikipedia.org/wiki/Bipartite_graph) and the weight of each edge is the rating given by user to the movie



the Adjacency matrix

	1	2	3
1	0	4	0
2	0	0	3
3	2	0	0
4	0	1	0
5	0	0	3

you can construct this matrix like $A[i][j] = r_{ij}$ here i is user_id, j is movie_id and r_{ij} is rating given by user i to the movie j

- we will Apply SVD decomposition on the Adjacency matrix [link1](https://stackoverflow.com/a/31528944/4084039) (<https://stackoverflow.com/a/31528944/4084039>), [link2](https://machinelearningmastery.com/singular-value-decomposition-for-machine-learning/) (<https://machinelearningmastery.com/singular-value-decomposition-for-machine-learning/>) and get three matrices U , Σ , V such that $U \times \Sigma \times V^T = A$, if A is of dimensions $N \times M$ then
 U is of $N \times k$,
 Σ is of $k \times k$ and
 V is $M \times k$ dimensions.
- So the matrix U can be represented as matrix representation of users, where each row u_i represents a k -dimensional vector for a user
- So the matrix V can be represented as matrix representation of movies, where each row v_j represents a k -dimensional vector for a movie
- μ represents the mean of all the rating given in the dataset

8.

for each epoch:

for each pair of (user, movie):

$b_i = b_i - \text{learning_rate} * dL/db_i$

$c_j = c_j - \text{learning_rate} * dL/dc_j$

predict the ratings with formula

$$\hat{y}_{ij} = \mu + b_i + c_j + \text{dot_product}(u_i, v_j)$$

print the mean squared error with predicted ratings

- you can choose any learning rate and regularization term in the range 10^{-3} to 10^2
- bonus:** instead of using SVD decomposition you can learn the vectors u_i , v_j with the help of SGD algo similar to b_i and c_j

TASK: 2

As we know U is the learned matrix of user vectors, with its i -th row as the vector u_i for user i . Each row of U can be seen as a "feature vector" for a particular user.

The question we'd like to investigate is this: do our computed per-user features that are optimized for predicting movie ratings contain anything to do with gender?

The provided data file [user_info.csv](https://drive.google.com/open?id=1PHFdJh_4gIPiLH5Q4UErH8GK71hTrzIY) (https://drive.google.com/open?id=1PHFdJh_4gIPiLH5Q4UErH8GK71hTrzIY) contains an `is_male` column indicating which users in the dataset are male. Can you predict this signal given the features U ?

Note 1 : there is no train test split in the data, the goal of this assignment is to give an intuition about how to do matrix factorization with the help of SGD and application of truncated SVD. for better understanding of the collaborative filtering please check netflix case study.

Note 2 : Check if scaling of U , V matrices improve the metric

In [1]:

```
1 import pandas as pd
2 from tqdm import tqdm
3 import matplotlib.pyplot as plt
4 from sklearn.metrics import mean_squared_error
```

In [2]:

```
1 data = pd.read_csv('ratings_train.csv')
2 data.shape
```

Out[2]:

(89992, 3)

In [3]:

```
1 from sklearn.utils.extmath import randomized_svd
2 import numpy as np
3 matrix = np.random.random((20, 10))
4 U, Sigma, VT = randomized_svd(matrix, n_components=5, n_iter=5, random_state=None)
5 print(U.shape)
6 print(Sigma.shape)
7 print(VT.T.shape)
```

(20, 5)

(5,)

(10, 5)

In []:

```
1
```

TASK:1

Constructing Adjacency Matrix

In [4]:

```
1 data.head()
```

Out[4]:

	user_id	item_id	rating
0	772	36	3
1	471	228	5
2	641	401	4
3	312	98	4
4	58	504	5

In [5]:

```
1 n=max(np.unique(data.user_id))+1
2 m=max(np.unique(data.item_id))+1
3 A=np.zeros((n,m), dtype=int)
4 print(n,m)
5
```

943 1681

In [6]:

```
1 for i in data.values:
2     A[i[0]][i[1]]=i[2]
```

In [7]:

```
1 # Adjacency Matrix
2 A
```

Out[7]:

```
array([[5, 0, 4, ..., 0, 0, 0],
       [4, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 0, 0],
       ...,
       [0, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 0, 0],
       [0, 5, 0, ..., 0, 0, 0]])
```

Applying SVD on matrix A to obtain the U, sigma and VT matrices

In [8]:

```
1 from sklearn.utils.extmath import randomized_svd
2 import numpy as np
3
4 U, Sigma, VT = randomized_svd(A, n_components=200, n_iter=5, random_state=None)
5 print(U.shape)
6 print(Sigma.shape)
7 print(VT.T.shape)
8 V=VT.T
```

```
(943, 200)
(200,)
(1681, 200)
```

Initialising the bias vectors for users and movies randomly from normal distribution.

In [11]:

```
1 data.head()
```

Out[11]:

	user_id	item_id	rating
0	772	36	3
1	471	228	5
2	641	401	4
3	312	98	4
4	58	504	5

In [12]:

```
1 B_i=np.random.rand(943)
2 C_j=np.random.rand(1681)
```

Training to find best B_i and C_i

$$L = \min_{b,c,\{u_i\}_{i=1}^N,\{v_j\}_{j=1}^M} \alpha \left(\sum_j \sum_k v_{jk}^2 + \sum_i \sum_k u_{ik}^2 + \sum_i b_i^2 + \sum_j c_j^2 \right) + \sum_{i,j \in \mathcal{I}^{\text{train}}} (y_{ij} - \mu - b_i - c_j - u_i^T v_j)^2$$

Initializing variables

In [13]:

```
1 mu=data['rating'].mean()
2 lr=1e-2
3 alpha=1
4 epochs=10
```

Functions to Find Derivative w.r.t B_i

$$dL/dB_i = 2\alpha(b_i) - 2(y_{ij} - \mu - b_i - c_j - u_i^T v_j)$$

In [14]:

```
1 def dl_dBi(b_i,c_j,u_i,v_j,y_ij):  
2     temp=np.dot(u_i.T,v_j)  
3     return 2*alpha*b_i-2*(y_ij-mu-b_i-c_j-temp)
```

Functions to Find Derivative w.r.t C_j

$$dL/dC_j = 2\alpha(c_j) - 2(y_{ij} - \mu - b_i - c_j - u_i^T v_j)$$

In [15]:

```
1 def dl_dCj(b_i,c_j,u_i,v_j,y_ij):  
2     temp=np.dot(u_i.T,v_j)  
3     return 2*alpha*c_j-2*(y_ij-mu-b_i-c_j-temp)
```

Training


In [16]:

```
1 MSE=[]
2 mse_prev=1;
3 for k in (range(epochs)):
4     for i in data.values:
5         b_i=B_i[i[0]]
6         c_j=C_j[i[1]]
7         u_i=U[i[0]]
8         v_j=V[i[1]]
9         y_ij=i[2]
10
11         # getting the derivatives with respect to b_i and c_j
12         dl_dbi=dl_dBi(b_i,c_j,u_i,v_j,y_ij)
13         dl_dcj=dl_dBi(b_i,c_j,u_i,v_j,y_ij)
14
15         # Updating b_i,c_j
16         B_i[i[0]]=B_i[i[0]]-lr*dl_dbi
17         C_j[i[1]]=C_j[i[1]]-lr*dl_dcj
18
19
20     y_ij_hat=[]
21     for i in data.values:
22         b_i=B_i[i[0]]
23         c_j=C_j[i[1]]
24         u_i=U[i[0]]
25         v_j=V[i[1]]
26
27         temp=mu+b_i+c_j+np.dot(u_i.T,v_j)
28         y_ij_hat.append(temp)
29
30     mse=mean_squared_error(data.rating, y_ij_hat)
31     if((abs(mse_prev-mse)/mse_prev)<0.001) :
32         break
33     mse_prev=mse
34     print("Epoch: ",k,":",mse)
35     MSE.append(mse)
36
```

```
Epoch: 0 : 0.964933675091534
Epoch: 1 : 0.9181488268959203
Epoch: 2 : 0.9046211967294691
Epoch: 3 : 0.8995164852409829
Epoch: 4 : 0.8974396889983892
```

In [17]:

```
1 x_epoch=[i for i in range(k)]
2 plt.plot(x_epoch, MSE, label="MSE")
3 plt.legend()
4 plt.xlabel("Epochs")
5 plt.ylabel("MSE")
6 plt.title("Error Plot")
7 plt.grid()
8 plt.show()
```



Task 2:

In [18]:

```
1 data2= pd.read_csv('user_info.csv')
2 data2.shape
```

Out[18]:

(943, 4)

In [19]:

```
1 data2.head()
```

Out[19]:

	user_id	age	is_male	orig_user_id
0	0	24	1	1
1	1	53	0	2
2	2	23	1	3
3	3	24	1	4
4	4	33	0	5

In [20]:

```
1 X=U;  
2 y=data2.is_male.values
```

In [21]:

```
1 # Train/test split data to apply classification algorithm  
2 from sklearn.model_selection import train_test_split  
3 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_state=
```

Hyperparameter Tuning to find best C for Logistic Regression

In [22]:

```
1
2 from sklearn.model_selection import GridSearchCV
3 from sklearn.linear_model import LogisticRegression
4
5 parameters = {'C':[0.00001,0.0001,0.001,0.01,0.1,1,10,100,1000]}
6 lr1 = LogisticRegression(random_state=0,max_iter=500)
7 clf = GridSearchCV(lr1, parameters,scoring='f1')
8 clf.fit(X_train,y_train)
9 results=pd.DataFrame(clf.cv_results_)
10 results
```

Out[22]:

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_C	params	split0_test_
0	0.006249	0.007653	0.000000	0.000000	1e-05	{'C': 1e-05}	0.8
1	0.000000	0.000000	0.005426	0.006772	0.0001	{'C': 0.0001}	0.8
2	0.000000	0.000000	0.003126	0.006252	0.001	{'C': 0.001}	0.8
3	0.003129	0.006258	0.000000	0.000000	0.01	{'C': 0.01}	0.8
4	0.000000	0.000000	0.003129	0.006258	0.1	{'C': 0.1}	0.8
5	0.007550	0.006999	0.000000	0.000000	1	{'C': 1}	0.8
6	0.009311	0.008628	0.000601	0.000802	10	{'C': 10}	0.8
7	0.019264	0.004504	0.000000	0.000000	100	{'C': 100}	0.8
8	0.055422	0.009758	0.000409	0.000817	1000	{'C': 1000}	0.8

Best Hyperparameter for training

In [24]:

```
1 best_C=clf.best_params_['C']
2 best_C
```

Out[24]:

10

Training with best hyperparameter

In [25]:

```
1 lr_=LogisticRegression(C=best_C,max_iter=500).fit(X_train, y_train)
```

In [26]:

```
1 # Predicted Value
2 y_pred=lr_.predict(X_test)
```

In [28]:

```
1 # calculating Mean square error
2 mse_lr=mean_squared_error(y_test, y_pred)
3 print("Mean Squared Error: ",mse_lr)
```

Mean Squared Error: 0.2564102564102564

Printing Confusing matrix

In [29]:

```
1 from sklearn.metrics import confusion_matrix
2 print("Confusion Matrix:")
3 confusion_matrix(y_test, y_pred)
```

Confusion Matrix:

Out[29]:

```
array([[ 24,  65],
       [ 15, 208]], dtype=int64)
```

Conclusion

We can see that the User matrix featurisation that we obtained from SVD is performing quite well in predicting the gender of the user with MSE = 0.256

Applying SGD for recommendation system after Normalising U(User Matrix) and V(Item Matrix)

Normalizing User matrix

In [31]:

```
1 from sklearn.preprocessing import Normalizer
2 normalizer_U=Normalizer()
3 U=normalizer_U.fit_transform(U)
4 U
```

Out[31]:

```
array([[ 0.08640271,  0.01029192, -0.01634916, ..., -0.04015678,
        -0.02687022,  0.01747944],
       [ 0.04026194, -0.14428045,  0.1668799 , ..., -0.09938937,
        0.08556978, -0.15432167],
       [ 0.02027543, -0.09368377,  0.07466926, ...,  0.06048021,
        -0.00515134, -0.06537516],
       ...,
       [ 0.02867392, -0.100791 ,  0.02461513, ..., -0.00043567,
        -0.00970243, -0.02189203],
       [ 0.04592986,  0.00822705,  0.04787223, ...,  0.0975489 ,
        -0.09033398,  0.05530162],
       [ 0.07458974, -0.00484077, -0.10452043, ...,  0.01343355,
        0.12689668, -0.0761227 ]])
```

Normalizing Item Matrix

In [33]:

```
1 from sklearn.preprocessing import Normalizer
2 normalizer_V=Normalizer()
3 V=normalizer_V.fit_transform(V)
4 V
```

Out[33]:

```
array([[ 0.11176055, -0.11097293, -0.00605184, ...,  0.06328592,
        -0.03491599,  0.00677097],
       [ 0.07861628, -0.01522179, -0.13737885, ...,  0.00291477,
        0.07314408, -0.03416641],
       [ 0.05535981, -0.07051791, -0.03214025, ..., -0.08229548,
        -0.13219517,  0.12801568],
       ...,
       [ 0.00281111, -0.04214132,  0.04742856, ..., -0.12968693,
        -0.09453178,  0.03178162],
       [ 0.00804536,  0.00273229, -0.0097731 , ...,  0.16962289,
        0.02779956, -0.00923656],
       [ 0.00809868,  0.00660892, -0.00780998, ..., -0.22814241,
        0.09825382, -0.11468601]])
```

In [40]:

```
1 # Random Initialisation of Bais vectors for users and items
2 B_i=np.random.rand(943)
3 C_j=np.random.rand(1681)
```

Applying SGD after normalising features U and V


In [41]:

```
1 MSE=[]
2 mse_prev=1;
3 for k in (range(epochs)):
4     for i in data.values:
5         b_i=B_i[i[0]]
6         c_j=C_j[i[1]]
7         u_i=U[i[0]]
8         v_j=V[i[1]]
9         y_ij=i[2]
10
11         # getting the derivatives with respect to b_i and c_j
12         dl_dbi=dl_dBi(b_i,c_j,u_i,v_j,y_ij)
13         dl_dcj=dl_dBi(b_i,c_j,u_i,v_j,y_ij)
14
15         # Updating b_i,c_j
16         B_i[i[0]]=B_i[i[0]]-lr*dl_dbi
17         C_j[i[1]]=C_j[i[1]]-lr*dl_dcj
18
19
20     y_ij_hat=[]
21     for i in data.values:
22         b_i=B_i[i[0]]
23         c_j=C_j[i[1]]
24         u_i=U[i[0]]
25         v_j=V[i[1]]
26
27         temp=mu+b_i+c_j+np.dot(u_i.T,v_j)
28         y_ij_hat.append(temp)
29
30     mse=mean_squared_error(data.rating, y_ij_hat)
31     if((abs(mse_prev-mse)/mse_prev)<0.001) :
32         break
33     mse_prev=mse
34     print("Epoch: ",k,":",mse)
35     MSE.append(mse)
36
```

```
Epoch: 0 : 0.9472836901144063
Epoch: 1 : 0.8907128611299193
Epoch: 2 : 0.8745494807880269
Epoch: 3 : 0.8690011067412411
Epoch: 4 : 0.8673085153489173
```

In [43]:

```
1 x_epoch=[i for i in range(k)]
2 plt.plot(x_epoch, MSE, label="MSE")
3 plt.legend()
4 plt.xlabel("Epochs")
5 plt.ylabel("MSE")
6 plt.title("Error Plot")
7 plt.grid()
8 plt.show()
```



we see that after normalizing matrices U and V we get a slightly lower MSE

End